

td

# Inside ActiveJob

Jerry D'Antonio  
@jerrydantonio  
[jerrydantonio.com](http://jerrydantonio.com)

# About Me

- I live and work in Akron, OH
- I work at Test Double
  - Mission: Improve the way the world builds software
- I am the creator and lead maintainer of Concurrent Ruby
  - Used by Rails, Sidekiq, Elasticsearch, MS Azure, Sucker Punch, others
  - <http://concurrent-ruby.com>



# Assumptions

- You know what ActiveJob is
- You understand what tasks are best suited for asynchronous job processing
- You have used ActiveJob in a production Rails application
- You are familiar with at least one supported job processor
- You have a basic understanding of concurrency and parallelism

# What is ActiveJob?

“Active Job is a **framework** for declaring jobs and making them run on a variety of queueing backends. These jobs can be everything from regularly **scheduled** clean-ups, to billing charges, to mailings. Anything that can be chopped up into small units of work and run in **parallel**, really.”

--[http://guides.rubyonrails.org/active\\_job\\_basics.html](http://guides.rubyonrails.org/active_job_basics.html)

- Framework: It unifies pre-existing and subsequently-created job runners
  - Supports Backburner, Delayed Job, Qu, Que, Resque, Sidekiq, Sneakers, Sucker Punch
- Scheduled: Supports ASAP and time-based scheduling
- Parallel: Can potentially scale across processors and machines

# Why do we need ActiveJob?

- Background job processors exist because they solve a problem
  - Long-running tasks block your HTTP response
  - Not every operation must occur before you send a response
  - Some tasks can be performed after the response so long as there is a guarantee
  - Background job processors provide the scheduling and guarantees

# Why do we need ActiveJob?

- ActiveJob unifies the ecosystem of job processors
  - Job processors existed but each was unique
  - Switching from one to another often required significant refactoring
  - ActiveJob provides an abstraction layer which supports the most common features
  - “Picking your queuing backend becomes more of an operational concern.”
    - Rails Guides

# A Simple Job Class

```
class DoSomethingLaterJob < ActiveJob::Base

  queue_as :default # optional

  def perform(*args)

    # Do something later

  end

end
```

# Using Our Simple Job Class

```
# config/application.rb
```

```
module YourApp
  class Application < Rails::Application
    config.active_job.queue_adapter = :inside_job
  end
end
```

```
# later, probably in a controller
```

```
DoSomethingLaterJob.perform_later(user)
```

```
DoSomethingLaterJob.set(wait_until: Date.tomorrow.noon).perform_later(user)
```



# Building Our Async Backend

- Threaded or forked?
  - We're going to use concurrent-ruby thread pools
- Persisting our jobs
  - For simplicity, we'll store our job data in-memory (no database persistence)
  - Our job processor will only be suitable for development and testing

# Building Our Async Backend

- The pieces
  - `ActiveJob::Core`—the job metadata
  - Queue adapter—marshals the job between Rails and the job runner
  - Job runner—provides asynchronous behavior
- The job runner is independent of Rails (Sidekiq, Sucker Punch, etc.)
- The queue adapter is in the Rails ActiveJob gem

# ActiveJob::Core Class

- Is the object passed into the queue adapter when a job is enqueued
- Provides two hugely important methods:
  - **#serialize**
  - **#deserialize**

# ActiveJob::Core Class

- Has several useful attributes:
  - **:queue\_name**
  - **:priority**
  - **:scheduled\_at**
  - **:job\_id**
  - **:provider\_job\_id**

# Our Queue Adapter

```
class InsideJobAdapter
  def enqueue(job)
    InsideJob.enqueue(job.serialize, queue: job.queue_name)
  end

  def enqueue_at(job, timestamp)
    InsideJob.enqueue_at(job.serialize, timestamp, queue: job.queue_name)
  end
end
```

# Our Job Runner: The Thread Pool

- What we need
  - Jobs are post to queues
  - Jobs must run asynchronously
- What we have
  - Thread pools from Concurrent Ruby each have their own queue and one or more threads
  - So a new thread pool for each job queue is all we need

# Our Job Runner: Creating Queues

```
QUEUES = Concurrent::Map.new do |hash, queue_name|  
  hash.compute_if_absent(queue_name) do  
    InsideJob.create_thread_pool  
  
  end  
  
end  
  
...  
  
def create_thread_pool  
  Concurrent::CachedThreadPool.new  
  
end
```

## Our Job Runner: Enqueue Now

```
def enqueue(job_data, queue: 'default')  
  QUEUES[queue].post(job_data) do |job|  
    ActiveJob::Base.execute(job)  
  end  
  
end
```



# Our Job Runner: Enqueue For Later

```
def enqueue_at(job_data, timestamp, queue: 'default')
  delay = timestamp - Time.current.to_f

  if delay > 0
    Concurrent::ScheduledTask.execute(
      delay, args: [job_data], executor: QUEUES[queue]) do |job|

      ActiveJob::Base.execute(job)
    end
  else
    enqueue(job_data, queue: queue)
  end
end
```

```
class InsideJob

  QUEUES = ThreadSafe::Cache.new do |hash, queue_name|
    hash.compute_if_absent(queue_name) { ActiveSupport::InsideJob.create_thread_pool }
  end

  class << self
    def create_thread_pool
      Concurrent::CachedThreadPool.new
    end

    def enqueue(job_data, queue: 'default')
      QUEUES[queue].post(job_data) { |job| ActiveSupport::Base.execute(job) }
    end

    def enqueue_at(job_data, timestamp, queue: 'default')
      delay = timestamp - Time.current.to_f
      if delay > 0
        Concurrent::ScheduledTask.execute(delay, args: [job_data], executor: QUEUES[queue]) do |job|
          ActiveJob::Base.execute(job)
        end
      else
        enqueue(job_data, queue: queue)
      end
    end
  end
end
end
```

```
module ActiveJob
  module QueueAdapters

    class InsideJobAdapter

      def enqueue(job)
        ActiveJob::InsideJob.enqueue(job.serialize, queue: job.queue_name)
      end

      def enqueue_at(job, timestamp)
        ActiveJob::InsideJob.enqueue_at(job.serialize, timestamp, queue: job.queue_name)
      end
    end
  end
end
```

# Rails AsyncJob

This repository

Search

Pull requestsIssuesGist

+▼

rails / rails

Watch ▼2,192Unstar30,012Fork12,182

<> Code

! Issues425

**🔗 Pull requests474**

⚡ Pulse

📊 Graphs

Initial implementation of ActiveJob AsyncAdapter. #21257

Edit

Mergedkaspth merged 1 commit into rails:master from unknown repository on Aug 25, 2015

💬 Conversation65

🔑 Commits1

📄 Files changed13

+183 -4

jdantonio commented on Aug 16, 2015

Now that `activesupport` has a runtime dependency on `concurrent-ruby`, we can begin taking advantage of those tools in more ways. This PR creates a simple asynchronous ActiveJob adapter that posts jobs to a concurrent-ruby thread pool. Within the context of ActiveJob it provides functionality comparable to `sucker_punch`. Rails 5 users will now be able to create simple asynchronous jobs without installing additional gems simply by setting the new adapter:

```
# config/application.rb
module YourApp
  class Application < Rails::Application
    config.active_job.queue_adapter = :async
  end
end
```

A simple benchmark script which compares enqueue performance vs. `sucker_punch` can be found [here](#). Performance is comparable on both Ruby 2.2.2 and JRuby 9000.

If this PR is accepted I can add more features such as job prioritization and per-queue thread pools.

Labels

None yet

Milestone

No milestone

Assignee

No one assigned

Notifications

🔊 Unsubscribe

You're receiving notifications because you were mentioned.

8 participants

# Subsequent AsyncJob Refactoring

- Collapse into one file
- Rename stuff
- Assign a provider job id
- Use one thread pool for everything
  - Throttle the thread pool
  - Configure the thread pool

# Going Further

- Sucker Punch
  - Threaded, in-memory, asynchronous job processor
  - Uses Concurrent Ruby thread pools
  - Decorates every job to provide job statistics and greater job control
  - Includes advanced shutdown behavior and options
- Sidekiq
  - Threaded job processor using database persistence
  - Includes many advanced features

td

My name is Jerry D'Antonio

Please tweet me @jerrydantonio &

Say hello@testdouble.com